

Importing a database table as pandas dataframe and some EDA

We will now be importing some of the tables from the **stats** database into this notebook as pandas dataframes.

```
In [1]: from sqlalchemy import create_engine
import pymysql
import pandas as pd
```

```
In [2]: # 'mysql+pymysql://mysql_user:mysql_password@mysql_host/mysql_db'
db_connection_str = 'mysql+pymysql://root:qaz2wsx@localhost/stats'
db_connection = create_engine(db_connection_str)
```

```
In [3]: # import the "posts" table from the 'stats' database into pandas dataframe
posts = pd.read_sql('SELECT * FROM posts', con=db_connection)
```

```
In [4]: posts.head(3)
```

```
Out[4]:
```

| | Id | PostTypeId | AcceptedAnswerId | ParentId | Score | ViewCount | Body | OwnerUserId | OwnerDisplayName | LastEditorUserId | LastEditDate |
|---|----|------------|------------------|----------|-------|-----------|---|-------------|------------------|------------------|---------------------|
| 0 | 1 | 1 | 15.0 | NaN | 49 | 4896.0 | <p>How should I elicit prior distributions fro... | 8.0 | None | NaN | NaT |
| 1 | 2 | 1 | 59.0 | NaN | 34 | 33132.0 | <p>In many different statistical methods there... | 24.0 | None | NaN | 2010-08-07 17:56:45 |
| 2 | 3 | 1 | 5.0 | NaN | 71 | 6503.0 | <p>What are some valuable Statistical Analysis... | 18.0 | None | 183.0 | 2011-02-12 05:50:04 |

```
In [5]: posts.shape
```

```
Out[5]: (372715, 18)
```

The *posts* dataframe has 372,715 rows and 18 columns.

```
In [6]: posts.dtypes
```

```
Out[6]:
```

| | |
|------------------|----------------|
| Id | int64 |
| PostTypeId | int64 |
| AcceptedAnswerId | float64 |
| ParentId | float64 |
| Score | int64 |
| ViewCount | float64 |
| Body | object |
| OwnerUserId | float64 |
| OwnerDisplayName | object |
| LastEditorUserId | float64 |
| LastEditDate | datetime64[ns] |
| LastActivityDate | datetime64[ns] |
| Title | object |
| Tags | object |
| AnswerCount | float64 |
| CommentCount | int64 |
| FavoriteCount | float64 |
| CreationDate | datetime64[ns] |
| dtype: | object |

We now also know the datatypes for all the columns in the *posts* dataframe. We will convert the datatype for any column(s), if necessary.

As you can see above, the *posts* dataframe has more than 370,000 rows. In the future, we might work with datasets that have millions of rows.

When we have such huge datasets, often the EDA and manipulation and training models takes very long. Let's see if we can work with a smaller subset of this dataset and still get similar (*as close as possible to the ground truth*) results.

We will do a very small experiment for this above task. We will pull random subsamples from the *posts* dataframe and compare the distributions of **score** with the distribution of score of the entire dataset.

Distribution of Score

Score in this dataset is defined as (number of upvotes - number of downvotes) for the posts (*which includes questions, answers, etc*).

First the distribution of **score** for the entire dataset.

```
In [7]: posts['Score'].describe()
```

```
Out[7]: count    372715.000000
mean         3.166299
std          9.712236
min         -58.000000
25%          1.000000
50%          1.000000
75%          3.000000
max         1714.000000
Name: Score, dtype: float64
```

As we can see above, the mean score is 3.16, and the median is 1.0 -- this suggests that our distribution is right skewed.

Let's confirm this further.

```
In [8]: posts['Score'].quantile(0.9)
```

```
Out[8]: 6.0
```

```
In [9]: posts['Score'].quantile(0.95)
```

```
Out[9]: 11.0
```

As suspected, 95% of the posts have a score of less than equal to 11 (*when the maximum goes upto 1714*). Only 5% of the posts have high scores, which are responsible for a high mean, and an even higher standard deviation.

Let's check this hypothesis as well.

Let us create a new dataframe **posts_95percentile** containing only the rows that have a score of less than equal to 11.

```
In [10]: posts_95percentile = posts.loc[posts['Score'] <= 11]
```

Now, let's check the distribution once again.

```
In [11]: posts_95percentile['Score'].describe()
```

```
Out[11]: count    356220.000000
mean         2.014609
std          2.256631
min         -58.000000
25%          1.000000
50%          1.000000
75%          3.000000
max         11.000000
Name: Score, dtype: float64
```

The **mean** has dropped from 3.16 to 2.01 and **standard deviation** from 9.71 to 2.25 (*i.e. the score is less spread now*).

Let us also visualise both the distributions.

```
In [12]: import seaborn as sns
```

```

import matplotlib.pyplot as plt

sns.set(style="darkgrid")

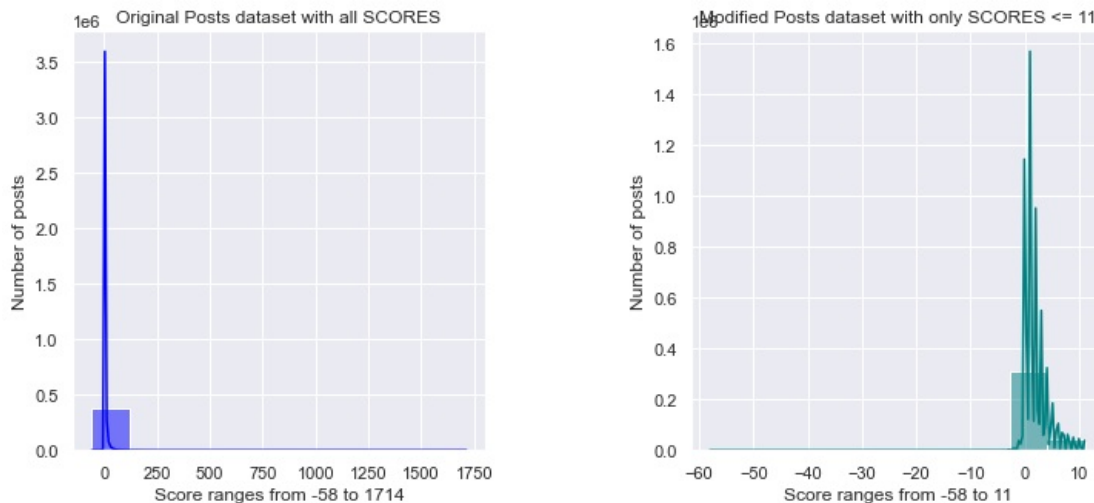
fig, (ax1,ax2) = plt.subplots(1, 2, figsize=(12, 5))

sns.histplot(data=posts, x="Score", kde=True, bins=10, color="blue", ax=ax1)
ax1.set_title('Original Posts dataset with all SCORES')
ax1.set_xlabel('Score ranges from -58 to 1714')
ax1.set_ylabel('Number of posts')

sns.histplot(data=posts_95percentile, x="Score", kde=True, bins=10, color="teal", ax=ax2)
ax2.set_title('Modified Posts dataset with only SCORES <= 11')
ax2.set_xlabel('Score ranges from -58 to 11')
ax2.set_ylabel('Number of posts')

fig.subplots_adjust(wspace=0.5)
plt.show()

```



As you might notice above, the distribution on the right seems a little bit more uniform than the one on the left.

You do not notice much variation in the sub-plots above because most of the posts don't get any upvotes or downvotes (hence, the peak around zero).

Nevertheless, we will now pull random subsamples from the **posts_95percentile** dataframe (*which has 356,220 rows*), and compare the distribution of scores.

Distribution of Score for different random subsamples

Let us pull three different subsamples from the dataframe **posts_95percentile** and compare their distribution of score(s). We will extract 25%, 50%, and 75% respectively of random data/rows from the dataframe.

```

In [13]: # pull random 25 percent of rows
posts_25percent_random = posts_95percentile.sample(frac=0.25)

# pull random 50 percent of rows
posts_50percent_random = posts_95percentile.sample(frac=0.50)

# pull random 75 percent of rows
posts_75percent_random = posts_95percentile.sample(frac=0.75)

```

```

In [14]: posts_25percent_random['Score'].describe()

```

```

Out[14]: count    89055.000000
mean         2.015451
std          2.270530
min         -58.000000
25%          1.000000
50%          1.000000
75%          3.000000
max          11.000000
Name: Score, dtype: float64

```

```

In [15]: posts_50percent_random['Score'].describe()

```

```

Out[15]: count    178110.000000

```

```

mean      2.016321
std       2.258812
min       -58.000000
25%       1.000000
50%       1.000000
75%       3.000000
max       11.000000
Name: Score, dtype: float64

```

```
In [16]: posts_75percent_random['Score'].describe()
```

```

Out[16]: count      267165.000000
mean       2.013220
std        2.254812
min        -58.000000
25%        1.000000
50%        1.000000
75%        3.000000
max        11.000000
Name: Score, dtype: float64

```

We do notice some subtle differences between the three subsamples. For example, the **posts_25percent_random** has a mean of 2.00 as compared to 2.01 for the original dataset (*very minute difference!*). For the same dataframe, the score at 25 percentile is still 0. The **posts_75percent_random** seems to be closest to the original dataset (*which isn't a surprise!*).

Having said the above, there isn't much difference between the three subsamples. This is primarily because most of the posts have a score of zero (as mentioned earlier). However, this might not be the case for future datasets and therefore, we have to be very careful when selecting a smaller subset especially, when training models (*high variance*).

Let us once visualise these subsamples as well.

```

In [17]: sns.set(style="darkgrid")

fig, axs = plt.subplots(2, 2, figsize=(8, 8))

sns.histplot(data=posts_95percentile, x="Score", kde=True, bins=10, color="teal", ax=axs[0,0])
axs[0,0].set_title('Original dataset')
axs[0,0].set_ylabel('Number of posts')

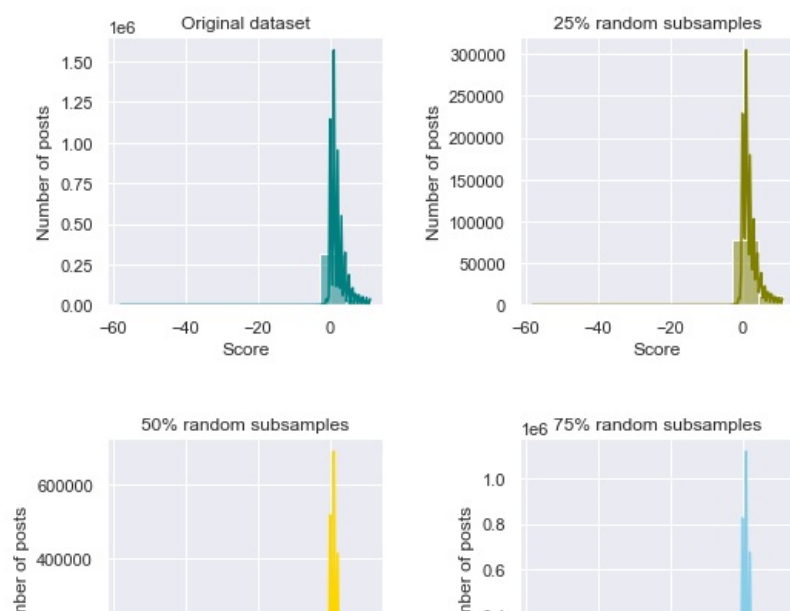
sns.histplot(data=posts_25percent_random, x="Score", kde=True, bins=10, color="olive", ax=axs[0,1])
axs[0,1].set_title('25% random subsamples')
axs[0,1].set_ylabel('Number of posts')


sns.histplot(data=posts_50percent_random, x="Score", kde=True, bins=10, color="gold", ax=axs[1,0])
axs[1,0].set_title('50% random subsamples')
axs[1,0].set_ylabel('Number of posts')

sns.histplot(data=posts_75percent_random, x="Score", kde=True, bins=10, color="skyblue", ax=axs[1,1])
axs[1,1].set_title('75% random subsamples')
axs[1,1].set_ylabel('Number of posts')

fig.subplots_adjust(wspace=0.5)
fig.subplots_adjust(hspace=0.5)
plt.show()

```





As mentioned above, not much difference between the three subsamples, and they all look very similar to the original histogram.

Let us explore our original **stats** database further.

This time let us also import the *tags* table into a dataframe, and see which topics get the most questions. We will also further explore how the number of questions have changed (increased/decreased) over the length of time.

Analysis on Tags

We will begin with importing the *tags* table into a pandas dataframe.

```
In [18]: tags = pd.read_sql('SELECT * FROM tags', con=db_connection)
```

```
In [19]: tags.head()
```

```
Out[19]:
```

| | Id | TagName | Count | ExcerptPostId | WikiPostId |
|---|----|---------------|-------|---------------|------------|
| 0 | 1 | bayesian | 6961 | 20258.0 | 20257.0 |
| 1 | 2 | prior | 866 | 62158.0 | 62157.0 |
| 2 | 3 | elicitation | 11 | NaN | NaN |
| 3 | 5 | open-source | 16 | NaN | NaN |
| 4 | 6 | distributions | 8428 | 8046.0 | 8045.0 |

TagName is the name of the tag and **Count** is the number of times the particular tag appears in the asked questions.

Let us observe the most used top 15 tags.

```
In [20]: tags_mostUsed = tags.sort_values(by = 'Count', ascending = False)
tags_top15 = tags_mostUsed.head(15)
```

```
In [21]: tags_top15
tags_top15.style.set_properties(**{'background-color': 'pink'}, subset=['Count'])
```

```
Out[21]:
```

| | Id | TagName | Count | ExcerptPostId | WikiPostId |
|-----|------|--------------------------|-------|---------------|--------------|
| 23 | 41 | r | 26385 | 2331.000000 | 2254.000000 |
| 49 | 111 | regression | 25540 | 8172.000000 | 8171.000000 |
| 5 | 9 | machine-learning | 17709 | 9066.000000 | 9065.000000 |
| 18 | 30 | time-series | 12763 | 3017.000000 | 3016.000000 |
| 86 | 188 | probability | 10758 | 20256.000000 | 20255.000000 |
| 21 | 38 | hypothesis-testing | 9004 | 3019.000000 | 3018.000000 |
| 4 | 6 | distributions | 8428 | 8046.000000 | 8045.000000 |
| 665 | 1483 | self-study | 7517 | 51952.000000 | 51951.000000 |
| 35 | 62 | neural-networks | 7055 | 50760.000000 | 50759.000000 |
| 0 | 1 | bayesian | 6961 | 20258.000000 | 20257.000000 |
| 453 | 975 | logistic | 6895 | 28327.000000 | 28326.000000 |
| 96 | 213 | mathematical-statistics | 6659 | 33251.000000 | 33250.000000 |
| 47 | 94 | classification | 6175 | 20262.000000 | 20261.000000 |
| 26 | 47 | correlation | 5623 | 8556.000000 | 8555.000000 |
| 65 | 135 | statistical-significance | 5532 | 20260.000000 | 20259.000000 |

Let's also visualise the above.

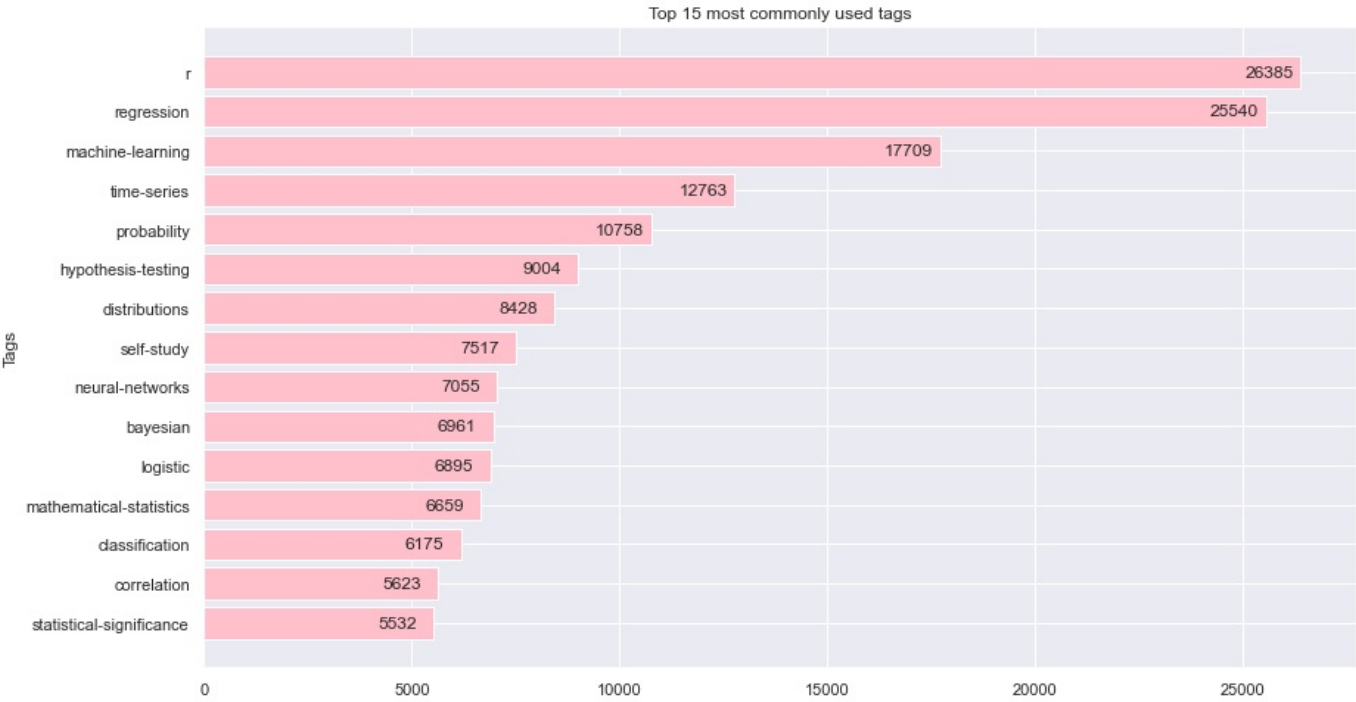
```
In [22]: tags_top15 = tags_top15.sort_values(by=['Count']) # to present the horizontal bars from most to least
```

```
In [23]: fig, ax = plt.subplots(figsize=(14, 8))

plt.barh(width = tags_top15.Count, y = tags_top15.TagName, color = 'pink')
plt.title('Top 15 most commonly used tags')
plt.ylabel('Tags')

ax.bar_label(ax.containers[0], padding = -38)

plt.show()
```



To be honest, the tag `r` is a surprise for me. Maybe, a lot of the users of stats.stackexchange.com use `r` (as the programming language) for their statistical analysis. This is followed by `regression`, `machine-learning`, `time-series`, and so on... (*self-study is an interesting choice for tag*)

Let us now see how the number of question have changed over time.

Number of questions over time

This time lets visualise how the number of questions posted on stats.stackexchange.com have changed in the past decade (*my money is on they must have increased as more and more people have increasingly gotten interested in data science/statistics*).

We will once again use the `posts` dataframe.

```
In [24]: posts.head(5)
```

| Out[24]: | Id | PostTypeId | AcceptedAnswerId | ParentId | Score | ViewCount | Body | OwnerUserId | OwnerDisplayName | LastEditorUserId | La |
|----------|----|------------|------------------|----------|-------|-----------|---|-------------|------------------|------------------|----|
| 0 | 1 | 1 | 15.0 | NaN | 49 | 4896.0 | <p>How should I elicit prior distributions fro... | 8.0 | None | NaN | |
| 1 | 2 | 1 | 59.0 | NaN | 34 | 33132.0 | <p>In many different statistical methods there... | 24.0 | None | NaN | |
| 2 | 3 | 1 | 5.0 | NaN | 71 | 6503.0 | <p>What are some valuable Statistical Analysis... | 18.0 | None | 183.0 | |
| 3 | 4 | 1 | 135.0 | NaN | 23 | 42655.0 | <p>I have two groups of data. Each with a dif... | 23.0 | None | NaN | |
| 4 | 5 | 2 | NaN | 3.0 | 90 | NaN | <p>The R-project</p>

<p><a href="http... | 23.0 | None | 23.0 | |

As we know from [here](#), **PostTypeId = 1** are questions and **PostTypeId = 2** are answers. Since we are only interested in questions, we will simply extract them from the `posts` dataframe.

```
In [25]: posts_questions = posts.loc[posts['PostTypeId'] == 1]
```

```
In [26]: posts_questions.shape
```

```
Out[26]: (187996, 18)
```

We have 187,996 questions.

Next, we have the **CreationDate** column which tells us when the question was posted. Since we are only interested in the year, we will simply extract just the relevant columns from our dataframe.

```
In [28]: posts_questionsANDyear = pd.DataFrame()
```

```
In [29]: posts_questionsANDyear['Body'], posts_questionsANDyear['Year'] = posts_questions['Body'], posts_questions['CreationDate']
```

```
In [30]: posts_questionsANDyear.head()
```

```
Out[30]:
```

| | Body | Year |
|---|---|------|
| 0 | <p>How should I elicit prior distributions fro... | 2010 |
| 1 | <p>In many different statistical methods there... | 2010 |
| 2 | <p>What are some valuable Statistical Analysis... | 2010 |
| 3 | <p>I have two groups of data. Each with a dif... | 2010 |
| 5 | <p>Last year, I read a blog post from <a href=... | 2010 |

We will now use the **groupby** method and get the *count* of how many questions were posted in a given year. We also want to return the result as a dataframe so that it is easier to visualise later.

```
In [31]: questions_by_year = posts_questionsANDyear.groupby('Year').size().reset_index(name='Count')
```

```
In [32]: questions_by_year
```

```
Out[32]:
```

| | Year | Count |
|----|------|-------|
| 0 | 2009 | 3 |
| 1 | 2010 | 1478 |
| 2 | 2011 | 4951 |
| 3 | 2012 | 8375 |
| 4 | 2013 | 11836 |
| 5 | 2014 | 15961 |
| 6 | 2015 | 18740 |
| 7 | 2016 | 20949 |
| 8 | 2017 | 21468 |
| 9 | 2018 | 19968 |
| 10 | 2019 | 19800 |
| 11 | 2020 | 21696 |
| 12 | 2021 | 22771 |

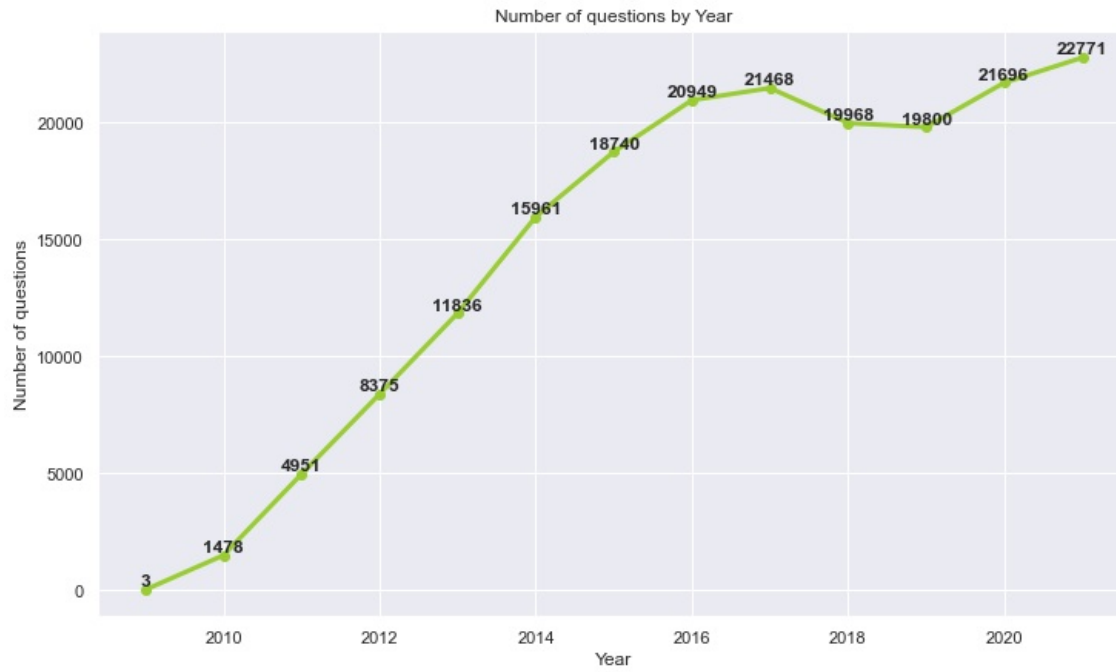
Let us also visualise the above.

```
In [51]: fig, ax = plt.subplots(figsize=(12, 7))
```

```
plt.plot('Year', 'Count', data=questions_by_year, marker='o', color='yellowgreen', linewidth=3)
plt.title('Number of questions by Year')
plt.ylabel('Number of questions')
plt.xlabel('Year')

for line in range(0, questions_by_year.shape[0]):
    plt.annotate(
        questions_by_year.Count[line],
        (questions_by_year.Year[line], questions_by_year.Count[line]+5),
        fontsize=12,
        weight='bold',
        va='bottom',
        ha='center'
    )

plt.show()
```



As predicted, the number of questions have generally increased with a slight dip in the years 2018 and 2019.

Question Tags over time

Now that we know that the number of questions being posted on stats.stackexchange.com have increased over time, let us also try to find out how the topics (*aka tags*) have evolved in the same time frame.

For this, we will create a new dataframe **posts_tagsANDyear** and extract only the relevant columns.

```
In [53]: posts_tagsANDyear = pd.DataFrame()
```

```
In [54]: posts_tagsANDyear['Body'], posts_tagsANDyear['Tags'], posts_tagsANDyear['Year'] = posts_questions['Body'], posts_questions['Tags'], posts_questions['Year']
```

```
In [56]: posts_tagsANDyear.head()
```

```
Out[56]:
```

| | Body | Tags | Year |
|---|---|---|------|
| 0 | <p>How should I elicit prior distributions fro... | <bayesian><prior><elicitation> | 2010 |
| 1 | <p>In many different statistical methods there... | <distributions><normality-assumption> | 2010 |
| 2 | <p>What are some valuable Statistical Analysis... | <software><open-source> | 2010 |
| 3 | <p>I have two groups of data. Each with a dif... | <distributions><statistical-significance> | 2010 |
| 5 | <p>Last year, I read a blog post from <a href=... | <machine-learning><pac-learning> | 2010 |

We see that there are multiple tags associated with any given question.

You might also remember that *Tags* data type is **object**.

Let's do some further data wrangling so that we have only one tag per row.

```
In [57]: import re

posts_tagsANDyear_modified = posts_tagsANDyear.assign(Tags=posts_tagsANDyear['Tags'].str.findall(r"(?<=<)(.*?)(?<=>|<|>|<br>|</br>|</p>|<p>|<a href=
```

```
In [58]: posts_tagsANDyear_modified.head(20)
```

Out[58]:

| | Body | Tags | Year |
|---|---|--------------------------|------|
| 0 | <p>How should I elicit prior distributions fro... | bayesian | 2010 |
| 0 | <p>How should I elicit prior distributions fro... | prior | 2010 |
| 0 | <p>How should I elicit prior distributions fro... | elicitation | 2010 |
| 1 | <p>In many different statistical methods there... | distributions | 2010 |
| 1 | <p>In many different statistical methods there... | normality-assumption | 2010 |
| 2 | <p>What are some valuable Statistical Analysis... | software | 2010 |
| 2 | <p>What are some valuable Statistical Analysis... | open-source | 2010 |
| 3 | <p>I have two groups of data. Each with a dif... | distributions | 2010 |
| 3 | <p>I have two groups of data. Each with a dif... | statistical-significance | 2010 |
| 5 | <p>Last year, I read a blog post from <a href=... | machine-learning | 2010 |
| 5 | <p>Last year, I read a blog post from <a href=... | pac-learning | 2010 |
| 6 | <p>I've been working on a new method for analy... | dataset | 2010 |
| 6 | <p>I've been working on a new method for analy... | sample | 2010 |
| 6 | <p>I've been working on a new method for analy... | population | 2010 |
| 6 | <p>I've been working on a new method for analy... | teaching | 2010 |
| 7 | <p>Sorry, but the emptyness was a bit overwhel... | humor | 2010 |
| 9 | <p>Many studies in the social sciences use Lik... | ordinal-data | 2010 |
| 9 | <p>Many studies in the social sciences use Lik... | likert | 2010 |
| 9 | <p>Many studies in the social sciences use Lik... | scales | 2010 |
| 9 | <p>Many studies in the social sciences use Lik... | measurement | 2010 |

This looks better. We will also *reset index*.

```
In [59]: posts_tagsANDyear_modified.reset_index(drop=True)
```

Out[59]:

| | Body | Tags | Year |
|--------|---|----------------------|------|
| 0 | <p>How should I elicit prior distributions fro... | bayesian | 2010 |
| 1 | <p>How should I elicit prior distributions fro... | prior | 2010 |
| 2 | <p>How should I elicit prior distributions fro... | elicitation | 2010 |
| 3 | <p>In many different statistical methods there... | distributions | 2010 |
| 4 | <p>In many different statistical methods there... | normality-assumption | 2010 |
| ... | ... | ... | ... |
| 575812 | <p>I'm working on a project which request to d... | predictive-models | 2021 |
| 575813 | <p>I'm working on a project which request to d... | model | 2021 |
| 575814 | <p>When I read up the example for Bayes Theore... | bayesian | 2021 |
| 575815 | <p>I want to train a model that takes image as... | neural-networks | 2021 |
| 575816 | <p>I want to train a model that takes image as... | conv-neural-network | 2021 |

575817 rows × 3 columns

```
In [86]: from wordcloud import WordCloud

def black_color_func(word, font_size, position, orientation, random_state=None, **kwargs):
    return("hsl(0,100%, 1%)")


tags_wc = WordCloud(font_path = './arial.ttf', max_words=30, background_color='white').generate(' '.join(posts_tagsANDyear_modified['Body']))
tags_wc.recolor(color_func = black_color_func)

<wordcloud.wordcloud.WordCloud at 0x1a1a2ff3e80>
```

Out [86]:

```
In [87]: plt.figure(figsize=[12,10])

plt.imshow(tags_wc, interpolation="bilinear")
plt.axis("off")
plt.show()
```



```
In [60]: tags_by_year = posts_tagsANDyear_modified.groupby(['Year', 'Tags'])['Year'].size().reset_index(name='Count')

tags_by_year
```

Out [60]:

| | Year | Tags | Count |
|-------|------|--------------------|-------|
| 0 | 2009 | cooks-distance | 1 |
| 1 | 2009 | data-visualization | 1 |
| 2 | 2009 | graph-theory | 1 |
| 3 | 2009 | moving-average | 1 |
| 4 | 2009 | outliers | 2 |
| ... | ... | ... | ... |
| 12843 | 2021 | z-score | 37 |
| 12844 | 2021 | z-statistic | 3 |
| 12845 | 2021 | z-test | 34 |
| 12846 | 2021 | zero-inflation | 49 |
| 12847 | 2021 | zipf | 3 |

12848 rows × 3 columns

References

1. <https://www.codegrepper.com/code-examples/sql/jupyter+notebook+mysql+connector+to+dataframe>
2. <https://stackoverflow.com/questions/37730243/importing-data-from-a-mysql-database-into-a-pandas-data-frame-including-column-n/37730334>
3. <https://python-graph-gallery.com/25-histogram-with-several-variables-seaborn>
4. <https://stackoverflow.com/questions/62351154/how-to-get-different-titles-for-each-of-the-subplots-in-seaborn>
5. <https://pandas.pydata.org/pandas-docs/stable/reference/api/pandas.DataFrame.sample.html>
6. <https://stackoverflow.com/questions/30228069/how-to-display-the-value-of-the-bar-on-each-bar-with-pyplot-barh>
7. <https://www.geeksforgeeks.org/how-to-annotate-matplotlib-scatter-plots/>
8. <https://www.analyticsvidhya.com/blog/2021/08/creating-customized-word-cloud-in-python/>